

# Titanic Survival Analysis and Prediction

## Author

DINESH S

[my linkedin profile](#) | [github link](#) | [Data Play](#)

## The Challenge

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

In this challenge, we ask you to build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socio-economic class, etc).



## Overview of Data

The data has been split into two groups:

- training set (train.csv)
- test set (test.csv)

The training set should be used to build your machine learning models.

The test set should be used to see how well your model performs on unseen data.

use the model you trained to predict whether or not they survived the sinking of the Titanic.

Data set Link: <https://www.kaggle.com/c/titanic/data>

importing the necessary libraries

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

```
In [ ]: #Default theme
sns.set_theme(context='notebook',
               style='whitegrid',
               palette='rainbow',
               font='Lucida Calligraphy',
               font_scale=1.5,
               rc=None)

import matplotlib
matplotlib.rcParams['figure.figsize'] = [8, 8]
matplotlib.rcParams.update({'font.size': 15})
matplotlib.rcParams['font.family'] = 'sans-serif'
```

Analyzing the Data set

```
In [ ]: titanic = pd.read_csv('/content/train.csv')
titanic.head().style.set_properties(
    **{
        'background-color': 'LightBlue',
        'color': 'Black',
        'border-color': 'darkblack'
    })
```

```
Out[ ]: 
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch
0	1	0	3	Braund, Mr. Owen Harris	male	22.000000	1	0
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38.000000	1	0
2	3	1	3	Heikkinen, Miss. Laina	female	26.000000	0	0
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.000000	1	0
4	5	0	3	Allen, Mr. William Henry	male	35.000000	0	0

```
In [ ]: print('Shape of Titanic data set is :',titanic.shape)
print('Size of Titanic data set is :',titanic.size)
```

```
Shape of Titanic data set is : (891, 12)
Size of Titanic data set is : 10692
```

```
In [ ]: titanic.describe()
```

Out[ ]:	PassengerId	Survived	Pclass	Age	SibSp	Parc
<b>count</b>	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000
<b>mean</b>	446.000000	0.383838	2.308642	29.699118	0.523008	0.38159
<b>std</b>	257.353842	0.486592	0.836071	14.526497	1.102743	0.80605
<b>min</b>	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000
<b>25%</b>	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000
<b>50%</b>	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000
<b>75%</b>	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000
<b>max</b>	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000

## Variable Notes

- pclass: A proxy for socio-economic status (SES)
  1. 1st = Upper
  2. 2nd = Middle
  3. 3rd = Lower
- age: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5
- sibsp: The dataset defines family relations in this way...
- Sibling: brother, sister, stepbrother, stepsister
- Spouse = husband, wife (mistresses and fiancés were ignored)
- parch: The dataset defines family relations in this way... Parent = mother, father

In [ ]: `titanic.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

## Data Visualization


```

In [ ]: matplotlib.rcParams.update({'font.size': 30})

titanic.dtypes.value_counts().plot.pie(explode=[0.1, 0.1, 0.1],
                                         autopct='%1.2f%%',
                                         shadow=True)

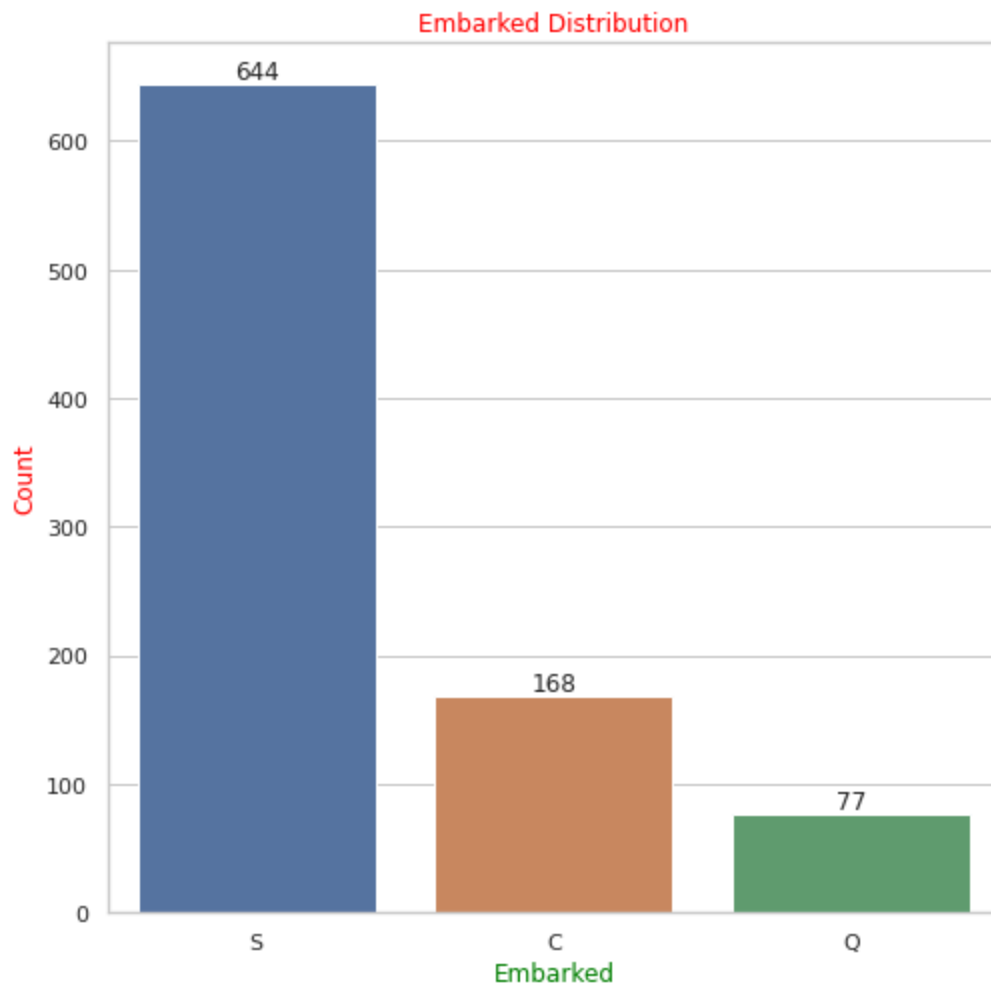
plt.title('Data Type',
          color='Green',
          loc='center',
          font='Lucida Calligraphy');

```




```
In [ ]: ax = sns.set(style="whitegrid")
ax = sns.countplot(data=titanic,x='Embarked');
ax.bar_label(ax.containers[0])

plt.title('Embarked Distribution',color='Red',loc='center',font='Lucida Call
plt.xlabel('Embarked',color='Green',loc='center',font='Lucida Calligraphy')
plt.ylabel('Count',color='Red',loc='center',font='Lucida Calligraphy');
```




```
In [ ]: matplotlib.rcParams.update({'font.size': 30})
titanic['Embarked'].value_counts().plot.pie(explode=[0.1, 0.1, 0.1],
                                              autopct='%1.2f%%',
                                              shadow=True)
plt.title('Embarked Distribution',color='Red',loc='center');
```

Embarked Distribution




```
In [ ]: matplotlib.rcParams.update({'font.size': 30})
titanic['Sex'].value_counts().plot.pie(explode=[0.1, 0.1],
                                         autopct='%1.2f%%',
                                         shadow=True)
plt.title('Gender Distribution', color='Red', loc='center');
```





```
In [ ]: matplotlib.rcParams.update({'font.size': 30})
titanic['Survived'].value_counts().plot.pie(explode=[0.1, 0.1],
                                              autopct='%1.2f%%',
                                              shadow=True)
plt.title('Survived', color='Red', loc='center');
```



```
In [ ]: titanic.corr().style.background_gradient(cmap='coolwarm').set_precision(3)
```

Out[ ]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
PassengerId	1.000	-0.005	-0.035	0.037	-0.058	-0.002	0.013
Survived	-0.005	1.000	-0.338	-0.077	-0.035	0.082	0.257
Pclass	-0.035	-0.338	1.000	-0.369	0.083	0.018	-0.549
Age	0.037	-0.077	-0.369	1.000	-0.308	-0.189	0.096
SibSp	-0.058	-0.035	0.083	-0.308	1.000	0.415	0.160
Parch	-0.002	0.082	0.018	-0.189	0.415	1.000	0.216
Fare	0.013	0.257	-0.549	0.096	0.160	0.216	1.000

```
In [ ]: corr=titanic.corr()["Survived"]
plt.figure(figsize=(20, 15))
sns.heatmap(corr, vmax=.8, linewidths=0.01, square=True, annot=True, cmap='YlGn')
plt.title('Correlation between features')
plt.show()
```




```
In [ ]: # correlation heatmap of highly correlated features with SalePrice
```


```
hig_corr = titanic.corr()
hig_corr_features = hig_corr.index[abs(hig_corr["Fare"]) >= 0.25]
hig_corr_features
```

```
Out[ ]: Index(['Survived', 'Pclass', 'Fare'], dtype='object')
```

```
In [ ]: plt.figure(figsize=(10,8))
ax = sns.heatmap(titanic[hig_corr_features].corr(), cmap = "coolwarm", annot
# to fix the bug "first and last row cut in half of heatmap plot"
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
plt.show()
```



```
In [ ]: g = sns.FacetGrid(titanic, col="Pclass")
g = g.map(plt.hist, "Age")
```



## Methods to find Missing Values

```
In [ ]: def missing_value (df):
    missing_Number = df.isnull().sum().sort_values(ascending=False)[df.isnull().sum().sort_values(ascending=False).index[0]]
    missing_percent=round((df.isnull().sum()/df.isnull().count()*100,2)[0],2)
    missing = pd.concat([missing_Number,missing_percent],axis=1,keys=['Missing Number','Missing Percent'])
    return missing
```


```
In [ ]: missing_value(titanic).style.background_gradient(cmap='coolwarm').set_precision(2)
```

Out[ ]:


	Missing Number	Missing Percentage
Cabin	687	77.10
Age	177	19.87
Embarked	2	0.22

```
In [ ]: missing_values = titanic.isnull().sum()
missing_values = missing_values[missing_values > 0]
missing_values.sort_values(inplace=True)
missing_values.plot.pie(explode=[0.1, 0.1, 0.1],
                        autopct='%1.2f%%',
                        shadow=True)


plt.title('Missing Values',
          color='Green',
          loc='center',
          font='Lucida Calligraphy');
```




```
In [ ]: import missingno
missingno.bar(titanic, color="dodgerblue", sort="ascending", figsize=(10,5),
```



```
In [ ]: missingno.matrix(titanic, figsize=(10,5), fontsize=12, color=(1, 0.38, 0.27))
```



```
In [ ]: missingno.heatmap(titanic, figsize=(10,5), fontsize=12);
```



# Feature Engineering

## Filling/Removing Missing Values

```
In [ ]: titanic['Age'] = titanic['Age'].fillna(titanic['Age'].mean())
```

```
In [ ]: titanic[titanic['Embarked'].isnull()]
```


```
Out[ ]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
<b>61</b>	62	1	1	Icard, Miss. Amelie	female	38.0	0	0	11357
<b>829</b>	830	1	1	Stone, Mrs. George Nelson (Martha Evelyn)	female	62.0	0	0	11357

```
In [ ]: titanic['Embarked'] = titanic['Embarked'].fillna(method='bfill')
```

```
In [ ]: titanic = titanic.drop(['Cabin'],axis=1)
```

```
In [ ]: import missingno as msno
msno.matrix(titanic)
plt.show()
```



```
In [ ]: titanic.isnull().sum()
```

```
Out[ ]: PassengerId    0
        Survived      0
        Pclass        0
        Name          0
        Sex           0
        Age           0
        SibSp         0
        Parch         0
        Ticket        0
        Fare          0
        Embarked      0
dtype: int64
```

## All the Missing Value is Filled/Removed

```
In [ ]: titanic = titanic.drop(['Name', 'Ticket'], axis=1)
```

```
In [ ]: titanic.head()
```

```
Out[ ]:   PassengerId  Survived  Pclass    Sex  Age  SibSp  Parch    Fare  Embark
0           1         0         3  male  22.0     1     0   7.2500
1           2         1         1 female  38.0     1     0  71.2833
2           3         1         3 female  26.0     0     0   7.9250
3           4         1         1 female  35.0     1     0  53.1000
4           5         0         3  male  35.0     0     0   8.0500
```

```
In [ ]: titanic = pd.get_dummies(titanic, columns=['Sex', 'Embarked'], drop_first=True)
        titanic.head()
```



```
Out[ ]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Emb
0	1	0	3	22.0	1	0	7.2500	1	
1	2	1	1	38.0	1	0	71.2833	0	
2	3	1	3	26.0	0	0	7.9250	0	
3	4	1	1	35.0	1	0	53.1000	0	
4	5	0	3	35.0	0	0	8.0500	1	

## Train Test Split

```
In [ ]: X = titanic.drop(['Survived'],axis=1)
y = titanic['Survived']
```

```
In [ ]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_st
```

## Standardizing the data

```
In [ ]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

X_train = pd.DataFrame(X_train, columns=X.columns)
X_test = pd.DataFrame(X_test, columns=X.columns)
```

```
In [ ]: display(X_train.head())
display(X_test.head())
```

	PassengerId	Pclass	Age	SibSp	Parch	Fare	Sex_male
0	1.360492	-1.584396	0.010681	-0.479698	-0.460682	-0.018600	0.728823
1	-1.632266	-1.584396	-0.119643	-0.479698	-0.460682	0.079245	0.728823
2	-1.344650	-1.584396	-0.503148	-0.479698	0.810657	0.646624	0.728823
3	-1.686680	-0.381742	-1.193456	0.493365	-0.460682	-0.031329	-1.372075
4	-1.111449	0.820913	0.033758	-0.479698	-0.460682	-0.479818	0.728823

	PassengerId	Pclass	Age	SibSp	Parch	Fare	Sex_male
0	0.676433	0.820913	-0.273045	0.493365	-0.460682	-0.315867	-1.372075
1	-0.248601	0.820913	-0.809952	-0.479698	-0.460682	-0.485419	0.728823
2	1.096196	0.820913	-0.733251	-0.479698	-0.460682	-0.467343	0.728823
3	1.488753	0.820913	0.010681	-0.479698	-0.460682	0.506858	0.728823
4	0.027354	-0.381742	0.493964	0.493365	2.081997	-0.078596	0.728823

# Model Implementation

## LogisticRegression

```
In [ ]: from sklearn.metrics import accuracy_score
# Logistic Regression
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
Y_pred = logreg.predict(X_test)

log_train = round(logreg.score(X_train, y_train) * 100, 2)
log_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :", log_train)
print("Model Accuracy Score :", log_accuracy)
```

Training Accuracy : 80.2  
Model Accuracy Score : 79.89

## Support Vector Machines

```
In [ ]: # Support Vector Machines
from sklearn.svm import SVC
svc = SVC()
svc.fit(X_train, y_train)
Y_pred = svc.predict(X_test)

svc_train = round(svc.score(X_train, y_train) * 100, 2)
svc_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :", svc_train)
print("Model Accuracy Score :", svc_accuracy)
```

Training Accuracy : 85.11  
Model Accuracy Score : 80.45

## KNeighborsClassifier

```
In [ ]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, y_train)
Y_pred = knn.predict(X_test)

knn_train = round(knn.score(X_train, y_train) * 100, 2)
knn_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :",knn_train)
print("Model Accuracy Score  :",knn_accuracy)
```

Training Accuracy : 90.03  
Model Accuracy Score : 75.98

## GaussianNB

```
In [ ]: # Gaussian Naive Bayes
from sklearn.naive_bayes import GaussianNB
gaussian = GaussianNB()
gaussian.fit(X_train, y_train)
Y_pred = gaussian.predict(X_test)

gaussian_train = round(gaussian.score(X_train, y_train) * 100, 2)
gaussian_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :",gaussian_train)
print("Model Accuracy Score  :",gaussian_accuracy)
```

Training Accuracy : 79.21  
Model Accuracy Score : 81.56

## LinearSVC

```
In [ ]: # Linear SVC
from sklearn.svm import LinearSVC
linear_svc = LinearSVC()
linear_svc.fit(X_train, y_train)
Y_pred = linear_svc.predict(X_test)

linear_svc_train = round(linear_svc.score(X_train, y_train) * 100, 2)
linear_svc_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :",linear_svc_train)
print("Model Accuracy Score  :",linear_svc_accuracy)
```

Training Accuracy : 80.34  
Model Accuracy Score : 81.01

## DecisionTreeClassifier

```
In [ ]: # Decision Tree
from sklearn.tree import DecisionTreeClassifier
decision = DecisionTreeClassifier()
decision.fit(X_train, y_train)
Y_pred = decision.predict(X_test)

decision_train = round(decision.score(X_train, y_train) * 100, 2)
decision_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :",decision_train)
print("Model Accuracy Score  :",decision_accuracy)
```

Training Accuracy : 100.0  
Model Accuracy Score : 72.63

## RandomForestClassifier

```
In [ ]: # Random Forest
from sklearn.ensemble import RandomForestClassifier
random_forest = RandomForestClassifier(n_estimators=100)
random_forest.fit(X_train, y_train)
Y_pred = random_forest.predict(X_test)
random_forest.score(X_train, y_train)

random_forest_train = round(random_forest.score(X_train, y_train) * 100, 2)
random_forest_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :",random_forest_train)
print("Model Accuracy Score  :",random_forest_accuracy)
```

Training Accuracy : 100.0  
Model Accuracy Score : 81.56

## XGBClassifier

```
In [ ]: import xgboost as Xgb
xgb = Xgb.XGBClassifier()
xgb.fit(X_train,y_train)
Y_pred = xgb.predict(X_test)
xgb.score(X_train, y_train)

xgb_train = round(xgb.score(X_train, y_train) * 100, 2)
xgb_accuracy = round(accuracy_score(Y_pred, y_test) * 100, 2)

print("Training Accuracy      :",xgb_train)
print("Model Accuracy Score  :",xgb_accuracy)
```

Training Accuracy : 100.0  
Model Accuracy Score : 77.65

## Comparing Models

```
In [ ]: models = pd.DataFrame({
    'Model': [
        'Support Vector Machines', 'KNN', 'Logistic Regression',
        'Random Forest', 'Perceptron',
        'Stochastic Gradient Decent', 'Linear SVC', 'Decision Tree',
        'GaussianNB', 'MLPClassifier', 'XGBClassifier'
    ],
    'Training Accuracy': [
        log_train, svc_train, knn_train, gaussian_train, perceptron_train,
        linear_svc_train, sgd_train, decision_train, random_forest_train,
        mlp_train, xgb_train
    ],
    'Model Accuracy Score': [
        log_accuracy, svc_accuracy, knn_accuracy, gaussian_accuracy, percept
        linear_svc_accuracy, sgd_accuracy, decision_accuracy, random_forest_
        mlp_accuracy, xgb_accuracy
    ]
})
```

```
In [ ]: models.sort_values(by='Training Accuracy', ascending=False)
```

```
Out[ ]:
```

	Model	Training Accuracy	Model Accuracy Score
7	Decision Tree	100.00	72.63
8	GaussianNB	100.00	81.56
10	XGBClassifier	100.00	77.65
2	Logistic Regression	90.03	75.98
9	MLPClassifier	86.38	79.89
1	KNN	85.11	80.45
5	Stochastic Gradient Decent	80.34	81.01
0	Support Vector Machines	80.20	79.89
6	Linear SVC	79.78	81.01
3	Random Forest	79.21	81.56
4	Perceptron	73.17	72.07

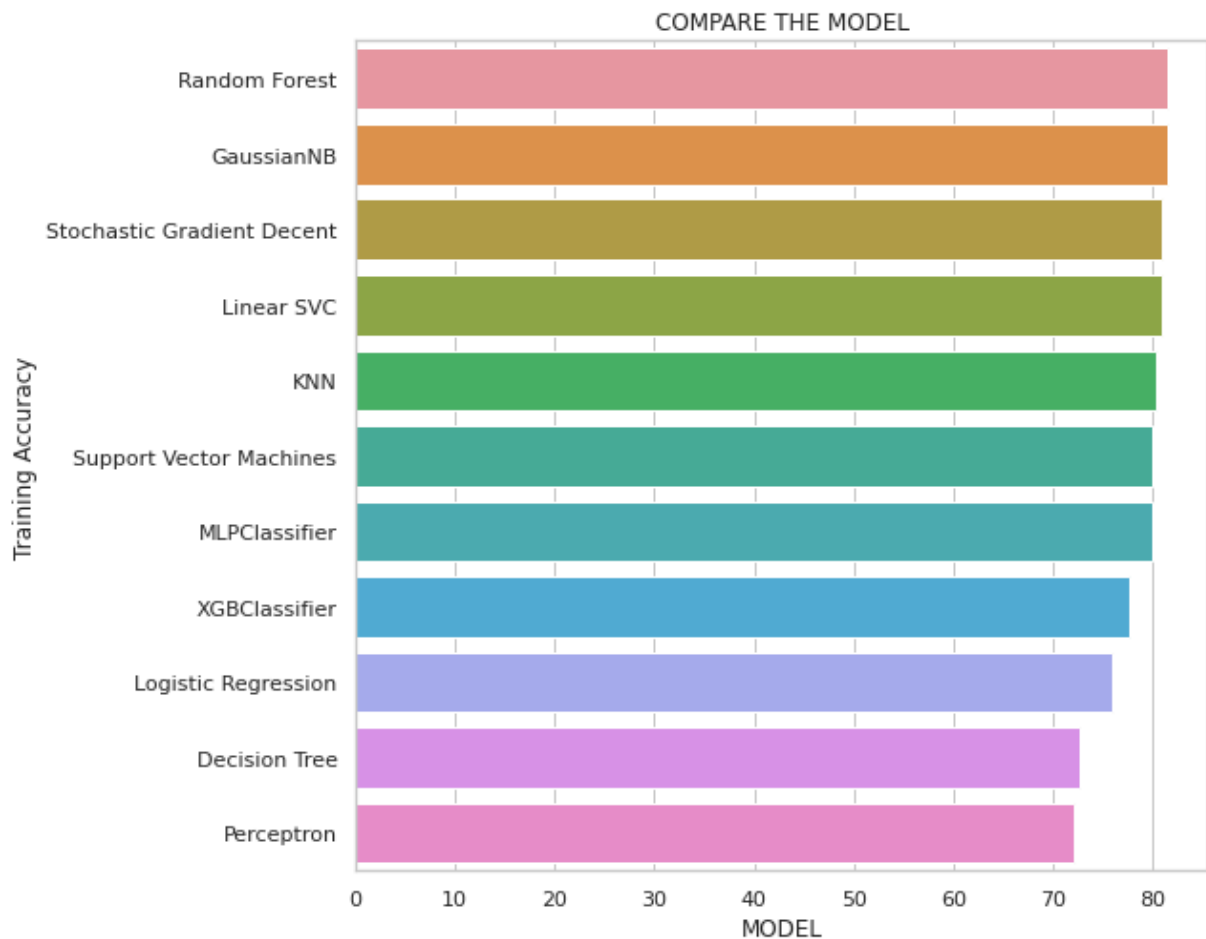
```
In [ ]: models.sort_values(by='Model Accuracy Score', ascending=False).style.background
    cmap='coolwarm').hide_index().set_properties(**{
        'font-family': 'Lucida Calligraphy',
        'color': 'LightGreen',
        'font-size': '15px'
    })
```

Out[ ]:

Model	Training Accuracy	Model Accuracy Score
Random Forest	79.210000	81.560000
GaussianNB	100.000000	81.560000
Stochastic Gradient Decent	80.340000	81.010000
Linear SVC	79.780000	81.010000
KNN	85.110000	80.450000
Support Vector Machines	80.200000	79.890000
MLPClassifier	86.380000	79.890000
XGBClassifier	100.000000	77.650000
Logistic Regression	90.030000	75.980000
Decision Tree	100.000000	72.630000
Perceptron	73.170000	72.070000

```
In [ ]: models=models.sort_values(by='Model Accuracy Score',ascending=False)[:20]

sns.barplot(y= 'Model', x= 'Model Accuracy Score', data= models)
plt.title('COMPARE THE MODEL')
plt.xlabel('MODEL')
plt.ylabel('Training Accuracy');
```



## Conclusion

In this notebook, we compared multiple machine learning models on both their training and test accuracy scores. Several observations stand out:

1. **Overfitting Indicators:** Models like Decision Tree, GaussianNB, and XGBClassifier achieved 100% training accuracy, but their test accuracies dropped—most notably, the Decision Tree's test accuracy fell to 72.63%, indicating overfitting.
2. **Top Performers:** Despite having perfect training scores, **GaussianNB** and **Random Forest** both reached the highest test accuracy of **81.56%**, suggesting they generalize well to unseen data compared to other models.
3. **Competitive Alternatives:** Stochastic Gradient Descent, Linear SVC, and KNN also performed relatively well, with test accuracies around 80–81%. Logistic Regression and MLPClassifier hovered in the mid-70s to high-70s range.
4. **Next Steps:**

- **Hyperparameter Tuning:** Fine-tune the top models (GaussianNB, Random Forest, etc.) to see if performance can be further improved.
- **Cross-Validation:** Employ more robust evaluation methods (e.g., k-fold cross-validation) for reliable performance estimates.
- **Additional Metrics:** Consider metrics like precision, recall, and F1-score to gain deeper insights into model performance.


Overall, GaussianNB and Random Forest emerged as strong contenders for this dataset. Further refinement and additional validation will help confirm which model offers the best balance of accuracy, efficiency, and interpretability.

## Acknowledgements

### Special Thanks:

I would like to extend my heartfelt gratitude to DataPlay Company for the fellowship. This opportunity has been instrumental in enhancing my skills and enabling projects like this to flourish.





---

*End of Notebook*

This notebook was converted with [convert.ploomber.io](https://convert.ploomber.io)